

UNIT 1: INTRODUCTION TO AI AND INTELLIGENT AGENTS

Course Pattern: TE IT - 2024 Pattern (SPPU)	Academic Term: Semester V (2026)
Status: Publication-Grade Theory Guide	Course Code: PCC-302-IT

Syllabus Mapping & Unit Structure

The following table outlines the exam-oriented curriculum mapping and priority weightage for Unit 1.

Syllabus Topic	In-Depth Exam Coverage	Priority / Weightage
Intro to AI & Turing Test	Definitions of AI (4 quadrants), Foundations of AI, Turing Test & Total Turing Test requirements.	High Priority (6-8 Marks Qs)
History & Types of AI	Evolution eras (Dartmouth 1956 to LLMs), Narrow AI (ANI) vs General AI (AGI) vs Super AI (ASI).	Medium Priority (4-6 Marks Qs)
AI vs ML vs DL vs DS	Hierarchical relationships, feature engineering comparison, detailed parameter matrix.	High Priority (6-8 Marks Qs)
Intelligent Agents & Rationality	Sensors, Actuators, Percept sequence P^* , Agent Function $f : P^* \rightarrow A$, Agent Program, Rationality vs Omniscience.	High Priority (7-9 Marks Qs)
PEAS & Environment Types	PEAS framework for Taxi, Medical, Tutor, Spam, Vacuum. 7 Environment dimensions.	High Priority (8-10 Marks Qs)
Types of Agent Architectures	Simple Reflex, Model-Based, Goal-Based, Utility-Based, Learning Agents (Critic, Learning/Performance Elements).	High Priority (8-10 Marks Qs)
Applications & AI Ethics	Applications in Healthcare, Education, Finance, Agriculture, Industry 4.0. Bias, XAI, Alignment, Governance.	Medium Priority (5-6 Marks Qs)

Introduction to Artificial Intelligence and AIML

Definitions of Artificial Intelligence

Artificial Intelligence (AI) is a multidisciplinary branch of computer science concerned with building smart machines capable of performing tasks that typically require human intelligence. Historically, researchers have categorized definitions of AI along two primary axes: **thought processes vs. behavior**, and **human performance vs. ideal performance (rationality)**.

Russell and Norvig organize these definitions into four distinct quadrants:

Human-Centric Focus	Rationality-Centric Focus
Thinking Humanly "The exciting new effort to make computers think... ma-	Thinking Rationally "The study of mental faculties through the use of com-

Human-Centric Focus	Rationality-Centric Focus
<i>chines with minds, in the full and literal sense.” (Haugeland, 1985)</i> Analyzes cognitive processing and internal thought mechanisms using cognitive science models.	<i>putational models.” (Charniak & McDermott, 1985)</i> Focuses on valid deductions, logic, syllogisms, and formal laws of thought.
Acting Humanly <i>“The art of creating machines that perform functions that require intelligence when performed by people.” (Kurzweil, 1990)</i> Evaluates operational output against human behavior (e.g., Turing Test benchmark).	Acting Rationally <i>“Computational Intelligence is the study of the design of intelligent agents.” (Poole et al., 1998)</i> Focuses on designing agents that operate to achieve the best outcome or best expected outcome.

The Turing Test and the Total Turing Test

Proposed by Alan Turing in his landmark 1950 paper “*Computing Machinery and Intelligence*”, the **Turing Test** (originally termed the Imitation Game) was designed to provide an operational definition of intelligence.

A computer passes the test if a human interrogator, after posing written questions, cannot reliably tell whether the written responses came from a person or a computer.

To pass the standard Turing Test, a machine must possess four core capabilities:

1. **Natural Language Processing (NLP):** To communicate successfully in human languages.
2. **Knowledge Representation:** To store what it knows or hears.
3. **Automated Reasoning:** To use stored information to answer questions and draw new conclusions.
4. **Machine Learning:** To adapt to new circumstances and detect patterns.

To pass the **Total Turing Test**, which includes physical signal interactions, the machine additionally requires:

5. **Computer Vision:** To perceive objects and physical environments.
6. **Robotics:** To manipulate objects and move around in the physical world.

Foundations of Artificial Intelligence

AI is built upon centuries of intellectual contributions from diverse disciplines:

- **Philosophy (410 BCE–present):** Formalized logic, theories of mind (empiricism vs. rationalism), and the connection between knowledge and physical action.
- **Mathematics (800 CE–present):** Provided formal logic (Boole, Frege), algorithms & computability (Turing, Gödel), and probability theory (Bayes, Bernoulli, Laplace) for reasoning under uncertainty.
- **Economics (1776–present):** Introduced decision theory, utility theory, and game theory for utility maximization under constraints.
- **Neuroscience (1861–present):** Examined brain anatomy and neural processing, inspiring artificial neural networks.
- **Psychology & Cognitive Science (1879–present):** Formulated theories of human memory, perception, and cognitive modeling.
- **Computer Engineering (1940–present):** Provided high-speed computational hardware, memory structures, and software frameworks needed to execute AI algorithms.
- **Control Theory & Cybernetics (1948–present):** Created self-regulating feedback loops and objective-driven state control mechanisms.
- **Linguistics (1957–present):** Formulated computational linguistics, generative grammars, and syntax analysis for natural language processing.

Source: Russell & Norvig 4th Ed, Ch 1, p. 19-34

History and Evolution of Artificial Intelligence

The history of AI spans several distinct phases marked by breakthroughs, high expectations, periods of reduced funding (AI Winters), and algorithmic revivals.

Era / Period	Key Milestones	Historical & Technical Significance
1943–1956: Gestation & Birth	McCulloch-Pitts Neuron (1943); Turing's paper (1950); Dartmouth Workshop (1956).	John McCarthy coined the term "Artificial Intelligence" at Dartmouth. Minsky, Rochester, and Shannon established AI as an independent field.
1952–1969: Early Enthusiasm	Newell & Simon's Logic Theorist & General Problem Solver (GPS); Samuel's Checkers; LISP language.	Demonstrated symbolic reasoning and heuristic search. Computers solved geometric theorems and algebraic word problems.
1966–1973: First AI Winter	Lighthill Report (1973); Minsky & Papert's <i>Perceptrons</i> book (1969).	Combinatorial explosion hindered progress. Proof that single-layer perceptrons could not solve linearly non-separable problems (XOR gate) halted neural network funding.
1969–1986: Expert Systems Boom	DENDRAL (molecular structure); MYCIN (blood infections); PROLOG & LISP machines.	Shifted focus from domain-independent search to domain-specific knowledge bases (Condition-Action production rules). Market grew to billions of dollars.
1987–1993: Second AI Winter	Collapse of specialized LISP machine hardware market; failure of Fifth Generation Computer Systems project.	Expert systems proved expensive to maintain, fragile, and difficult to update, leading to major government and corporate budget cuts.
1995–2010: Statistical ML & Big Data	IBM Deep Blue defeats Garry Kasparov (1997); Support Vector Machines (SVM); Probabilistic Graphical Models.	AI embraced rigorous mathematical and statistical methods, shifting from rule-based systems to data-driven probabilistic learning.
2012–Present: Deep Learning & GenAI	AlexNet wins ImageNet (2012); AlphaGo (2016); Transformer Architecture (2017); LLMs (ChatGPT, Gemini).	GPUs, deep neural networks, and massive internet-scale data enabled breakthrough perception, natural language understanding, and generative intelligence.

Source: Russell & Norvig 4th Ed, Ch 1, p. 35-49

Types of Artificial Intelligence

AI systems are categorized based on their capabilities, versatility, and degree of autonomy into three distinct evolutionary tiers:

1. Artificial Narrow Intelligence (ANI) / Weak AI

ANI refers to AI systems designed and trained to perform a specific, single task or a narrow range of defined tasks.

- **Characteristics:** Highly specialized, reactive, operating strictly within a pre-programmed or learned domain. ANI cannot transfer its intelligence or skills to an unlearned domain.
- **Current Status:** All existing operational AI systems today are ANI.
- **Examples:** Apple Siri, Google Translate, AlphaGo, recommendation engines (Netflix/Amazon), autonomous obstacle detection systems.

2. Artificial General Intelligence (AGI) / Strong AI

AGI refers to a theoretical AI system that possesses human-level cognitive capabilities across all operational domains.

- **Characteristics:** Ability to reason, plan, solve novel problems, abstract concepts, understand complex ideas, learn rapidly from experience, and transfer knowledge across unrelated domains.
- **Current Status:** Purely theoretical; actively researched under paradigms such as neuro-symbolic AI, cognitive architectures, and scalable transformer foundation models.

3. Artificial Super Intelligence (ASI)

ASI represents a hypothetical future stage where artificial intelligence surpasses human cognitive capabilities across every domain, including scientific creativity, general wisdom, social skills, and strategic planning.

- **Characteristics:** Self-improving recursive algorithms leading to an “intelligence explosion” (Technological Singularity).
- **Current Status:** Theoretical concept in AI safety and policy studies.

Parameter	ANI (Narrow AI)	AGI (General AI)	ASI (Super AI)
Scope	Single, specialized task	Multi-domain human equivalent	Exceeds all human intelligence
Adaptability	Fails outside trained domain	High domain transferability	Autonomous self-evolution
Reasoning	Statistical pattern matching	Abstract & common-sense reasoning	Hyper-dimensional problem solving
Current Status	Fully Deployed & Operational	Active Frontier Research	Theoretical / Conceptual

Source: Russell & Norvig 4th Ed, Ch 1, p. 49-53

AI vs Machine Learning vs Deep Learning vs Data Science

Understanding the operational boundaries and overlapping scopes of AI, Machine Learning (ML), Deep Learning (DL), and Data Science (DS) is essential.

Hierarchical Relationship

Mathematically and conceptually, these disciplines form a nested set structure:

$$\text{Deep Learning (DL)} \subset \text{Machine Learning (ML)} \subset \text{Artificial Intelligence (AI)}$$

Data Science is an interdisciplinary field that intersects with all three subsets while incorporating domain expertise, statistical analysis, and data engineering.

Disciplines Relationship Blueprint

ARTIFICIAL INTELLIGENCE (AI) — Broad field of smart, autonomous decision-making

MACHINE LEARNING (ML) — Learning patterns directly from data without explicit rules

DEEP LEARNING (DL) — Multi-layered neural network feature extraction

***DATA SCIENCE** intersects all layers, combining Big Data pipeline engineering, visualization, and domain modeling.*

Comparative Analysis Matrix

Parameter	Artificial Intelligence	Intelligence	Machine Learning	Deep Learning	Data Science
Primary Goal	Simulate human-like intelligence & decisions		Learn predictive models from data	Automatically extract representations via neural nets	Extract actionable business insights from data
Approach	Symbolic rules, logic, heuristics, search, ML		Statistical algorithms & mathematical optimization	Deep multi-layer Artificial Neural Networks (ANN)	Data cleaning, ETL, exploratory analysis, stats, ML
Feature Engineering	Hand-crafted rule systems & knowledge bases		Manual feature selection & engineering required	Automatic end-to-end representation learning	Feature extraction based on domain context
Data Requirement	Can work with explicit logic/rules (no data)		Medium-to-large structured datasets	Massive unstructured datasets (Text, Image, Video)	Works with structured, semi-structured & raw data
Hardware	Standard CPU / Microcontrollers		Standard CPU / Multi-core processors	Requires specialized GPUs / TPUs / Accelerators	Distributed compute clusters (Spark, Hadoop)
Interpretability	High (in rule-based expert systems)		Moderate (Decision Trees, Linear Models)	Low ("Black Box" non-linear representations)	High (Statistical metrics & visual reports)

Source: Alpaydin 3rd Ed, Ch 1, p. 1-15 | Source: Marsland 2nd Ed, Ch 1, p. 3-12

Intelligent Agents

Definition of an Agent

An **Agent** is anything that can be viewed as perceiving its environment through **sensors** and acting upon that environment through **actuators**.

- **Sensors:** Physical mechanisms (cameras, infrared range finders, sonar) or software inputs (keystrokes, network packets, file streams) that capture environment state.

- **Actuators:** Physical mechanisms (wheels, robotic arms, display screens) or software outputs (writing file packets, executing API calls) that alter the environment.
- **Percept:** The agent's perceptual inputs at any given instant.
- **Percept Sequence:** The complete history of everything the agent has ever perceived:

$$P^* = (p_1, p_2, p_3, \dots, p_t)$$

Agent Function vs. Agent Program

- **Agent Function (f):** An abstract mathematical function mapping any given percept sequence to an action:

$$f : P^* \rightarrow A$$

- **Agent Program:** The concrete implementation of the agent function running on a physical computational architecture:

$$\text{Agent} = \text{Architecture} + \text{Program}$$

EXAMPLE:

Vacuum-Cleaner World Percept Sequence & Action Mapping: Consider a simple two-location vacuum environment (Locations A and B).

- **Percepts:** [Location, Status] where Status is Clean or Dirty.
- **Actions:** Right, Left, Suck, No-Op.

Sample Agent Function $f(P^*)$ Table:

Percept Sequence P^*	Action A
([A, Clean])	Right
([A, Dirty])	Suck
([A, Clean], [B, Dirty])	Suck
([A, Clean], [B, Clean])	Left

The Concept of Rationality

A **Rational Agent** is one that selects an action that is expected to maximize its performance measure, given the evidence provided by the percept sequence and whatever built-in knowledge the agent has.

Rationality at any given time depends on four factors:

1. The **Performance Measure** that defines the criterion of success.
2. The agent's **Prior Knowledge** of the environment.
3. The **Actions** that the agent can perform.
4. The agent's **Percept Sequence** to date.

IMPORTANT:

Rationality vs. Omniscience:

- An *omniscient agent* knows the *actual* outcome of its actions and can act accordingly; omniscience is impossible in reality.
- A *rational agent* makes the *optimal decision based on expected outcomes* given its current information and bounded processing capabilities. Rationality encourages **information gathering, exploration, and learning** (autonomy).

Source: Russell & Norvig 4th Ed, Ch 2, p. 54-60

Agent Environment and PEAS Representation

The PEAS Framework

To design a rational agent, the task environment must be formally specified using the **PEAS** framework:

- **Performance Measure:** Quantitative metric evaluating agent success.
- **Environment:** External world/medium in which the agent operates.
- **Actuators:** Output devices/mechanisms used to execute actions.
- **Sensors:** Input devices/mechanisms used to perceive the environment.

Agent Type	Performance Measure	Environment	Actuators	Sensors
Automated Taxi Driver	Safety, speed, legal drive, profit, passenger comfort	City streets, traffic, pedestrians, weather, roadwork	Steering wheel, accelerator, brake, horn, display	Cameras, LiDAR, radar, GPS, speedometer, accelerometer
Medical Diagnostic System	Healthy patient, minimized costs, legal liability compliance	Patient, hospital staff, laboratory test feeds	Display screen (diagnosis, prescriptions, test orders)	Keyboard/touchscreen (symptom & lab data inputs)
Interactive English Tutor	Student test score maximization, engagement, speed	Student population, learning management platform	On-screen text, audio output, exercise recommendations	Keyboard, microphone, response time tracker
Spam Email Classifier	High precision/recall, zero false positive spam flags	Email server stream, user inbox, user flags	Move to spam folder, mark safe, flag user	Header parser, body text extractor, IP lookup
Vacuum Cleaner Agent	Cleanliness level, battery efficiency, time taken, low noise	Floor tiles, carpet, dirt, obstacles, walls	Drive wheels, suction motor, brush assembly	Dirt detection sensor, bump sensor, cliff sensor

Environment Properties (7 Dimensions)

The nature of the task environment directly dictates the complexity of the agent architecture:

1. Fully Observable vs. Partially Observable (vs. Unobservable):

- *Fully Observable:* Sensors detect the complete state of the environment at any given point in time (e.g., Chess).
- *Partially Observable:* Sensors have noisy, incomplete, or missing state information (e.g., Automated Taxi, Poker).

2. Single-Agent vs. Multi-Agent:

- *Single-Agent:* Agent operates alone without other decision-makers (e.g., Sudoku solver).
- *Multi-Agent:* Environment contains other agents acting competitively (Chess) or cooperatively (Taxi navigation).

3. Deterministic vs. Nondeterministic (Stochastic):

- *Deterministic:* The next environment state is completely determined by the current state and the executed action (e.g., Chess).
- *Nondeterministic / Stochastic:* Randomness or unknown external events make the next state uncertain (e.g., Driving in traffic).

4. Episodic vs. Sequential:

- *Episodic:* Agent's experience is divided into independent atomic episodes; actions in one episode do not affect future episodes (e.g., Defective part classification).
- *Sequential:* Current decisions affect all future decisions and states (e.g., Chess, Driving).

5. Static vs. Dynamic (vs. Semidynamic):

- *Static*: Environment does not change while the agent is deliberating (e.g., Crossword puzzle).
- *Dynamic*: Environment continuously changes while the agent thinks (e.g., Taxi driving).
- *Semidynamic*: Environment does not change with time, but the agent's performance score decreases with time (e.g., Timed chess).

6. Discrete vs. Continuous:

- *Discrete*: Environment states, time steps, perceptions, and actions are finite and distinct (e.g., Chess).
- *Continuous*: States and actions vary smoothly over continuous physical spectrums (e.g., Taxi steering angle, velocity).

7. Known vs. Unknown:

- *Known*: The physical/environmental rules governing outcomes are fully known to the agent.
- *Unknown*: The agent must learn the physics/rules of the environment to make decisions.

Source: Russell & Norvig 4th Ed, Ch 2, p. 60-67

Types of Agent Architectures

Agent programs are classified into five fundamental structural architectures based on their internal reasoning mechanisms:

1. Simple Reflex Agents

Simple Reflex Agents select actions based **only** on the current percept, ignoring the rest of the percept history. They operate using **Condition-Action Rules** (IF-THEN statements).

- **Mathematical Mapping**: $f(p) = \text{Rule-Action}(\text{Rule-Match}(\text{Current-State}))$
- **Limitations**: Only works if the environment is **fully observable**. In partially observable environments, infinite loops are common.

Simple Reflex Agent Architecture

Environment → Sensors → What the world is like now → Condition-Action Rules → What action I should do now → Actuators → Environment

2. Model-Based Reflex Agents

Model-Based Reflex Agents handle **partial observability** by maintaining an **internal state** that tracks unobserved aspects of the current world state.

• Required Models:

1. *Transition Model (How the world evolves)*: How the environment changes independently and in response to agent actions: $P(s_t \text{ mid } s_{t-1}, a_{t-1})$.
2. *Sensor Model (How the world reflects in percepts)*: How environment states translate to sensor readings: $P(e_t \text{ mid } s_t)$.

Model-Based Reflex Agent Architecture

Sensors → Internal State + How world evolves + What my actions do → What world is like now → Condition-Action Rules → Actuators

3. Goal-Based Agents

Goal-Based Agents combine internal state tracking with explicit **Goal Information** describing desirable situations.

- **Reasoning Mechanism:** The agent evaluates future sequences of actions to determine if they lead to the specified goal state (G). It integrates **Search** and **Planning** algorithms.
- **Advantage:** Flexible; if the goal changes (e.g., new navigation destination), the agent re-plans automatically without rewriting rule bases.

Goal-Based Agent Architecture

Sensors → State + World Evolution Model → What it will be like if I do Action X + Goals
→ Select Action achieving Goal → Actuators

4. Utility-Based Agents

Goal-based agents distinguish only between goal states and non-goal states. **Utility-Based Agents** use a continuous **Utility Function** ($U : S \rightarrow \mathbb{R}$) mapping environment states to real numbers, measuring the degree of “happiness” or desirability.

- **Decision Standard:** Maximizes **Expected Utility** under uncertainty:

$$\mathbb{E}[U(s')] = \sum_{s'} P(s' \text{ mid } s, a) U(s')$$

- **Advantages:** Trades off competing goals (e.g., speed vs. safety in driving) and handles stochastic environments rationally.

Utility-Based Agent Architecture

Sensors → State → What world will be like + Utility Function → Select Action Maximizing
Expected Utility → Actuators

5. Learning Agents

Learning Agents operate in initially unknown environments and improve their performance over time. A Learning Agent is divided into four conceptual components:

1. **Critic:** Evaluates the agent’s behavior against an external performance standard and provides feedback.
2. **Learning Element:** Responsible for making improvements based on feedback from the critic.
3. **Performance Element:** Responsible for selecting external actions (corresponds to the entire agent in previous architectures).
4. **Problem Generator:** Suggests exploratory actions that lead to new experiences rather than adhering strictly to known sub-optimal paths.

Learning Agent Structural Breakdown

Sensors → Critic (Feedback) → Learning Element (Learning Goals) → Problem Generator →
Performance Element → Actuators

Source: Russell & Norvig 4th Ed, Ch 2, p. 67-78

Applications of Artificial Intelligence in Key Domains

1. Healthcare

- **Computer-Aided Diagnosis (CAD):** Deep Convolutional Neural Networks (CNNs) detect malignant tumors in MRI, CT scans, and X-rays with expert-level precision.

- **Precision & Personalized Medicine:** Genomic data analysis models tailor drug treatment protocols based on individual patient genetic profiles.
- **Robotic Surgery:** Systems like the da Vinci Surgical System assist surgeons with micro-precision manipulations.

2. Education

- **Intelligent Tutoring Systems (ITS):** Platforms adapt difficulty and explanations based on real-time student performance metrics.
- **Automated Grading:** Natural Language Processing (NLP) models grade short answer essays and detect code plagiarism automatically.

3. Finance

- **Algorithmic High-Frequency Trading:** Predictive models execute stock trades at microsecond intervals based on market indicators.
- **Fraud Detection:** Anomaly detection models evaluate credit card transactions in real time to intercept unauthorized payments.

4. Agriculture

- **Precision Farming & Crop Health:** Computer vision systems mounted on drones detect weed infestations, nitrogen deficiencies, and soil moisture levels.
- **Autonomous Harvesters & Yield Prediction:** Computer vision-guided agricultural robots harvest delicate crops without damage.

5. Industry 4.0 & Smart Manufacturing

- **Predictive Maintenance:** IoT sensor streams coupled with time-series ML models forecast machine equipment failures before breakdowns occur.
- **Digital Twins:** Virtual simulation replicas of physical factories optimize assembly workflows and material handling.

Source: Russell & Norvig 4th Ed, Ch 1, p. 45-48

Challenges and Future of Artificial Intelligence

1. Ethical Considerations, Bias, and Fairness

- **Algorithmic Bias:** Machine learning models trained on historically biased data perpetuate discrimination in loan approvals, hiring tools, and judicial sentencing.
- **Data Privacy:** Massive data scraping for foundation models raises severe concerns regarding consent and personal information security.

2. Interpretability and Explainable AI (XAI)

- **The “Black Box” Problem:** Deep neural networks provide high accuracy but lack transparent decision pathways. XAI methods (e.g., LIME, SHAP) aim to make AI decisions interpretable in high-stakes domains such as healthcare and law.

3. The AI Alignment Problem

- **Goal Misalignment:** Ensuring that superintelligent or highly autonomous AI systems act in accordance with human ethical principles, safety guidelines, and intended outcomes rather than unintended proxy reward functions.

4. Economic & Societal Impact

- **Workforce Automation:** AI automation poses potential risks of workforce displacement while simultaneously creating demand for specialized data engineering and AI safety roles.

Source: Russell & Norvig 4th Ed, Ch 1, p. 49-53